

IMPORTING SQL

```
=====

from pyspark import SQLContext
sqlContext = SQLContext(sc)
departmentsJson = sqlContext.jsonFile("/user/cloudera/pyspark/departments.json")
departmentsJson.registerTempTable("departmentsTable")
departmentsData = sqlContext.sql("select * from departmentsTable")
for rec in departmentsData.collect():
    print(rec)

=====
```

SQL FROM PYSPARK

```
-----

from pyspark.sql import HiveContext
sqlContext = HiveContext(sc)
depts = sqlContext.sql("select * from departments")
for rec in depts.collect():
    print(rec)

sqlContext.sql("create table departmentsSpark as select * from departments")
depts = sqlContext.sql("select * from departmentsSpark")
for rec in depts.collect():
    print(rec)

=====
```

Using Hive

```
=====

from pyspark.sql import HiveContext
sqlContext = HiveContext(sc)
sqlContext.sql("set spark.sql.shuffle.partitions=10");

joinAggData = sqlContext.sql("select o.order_date,
round(sum(oi.order_item_subtotal), 2), \
count(distinct o.order_id) from orders o join order_items oi \
on o.order_id = oi.order_item_order_id \
group by o.order_date order by o.order_date")

for data in joinAggData.collect():
    print(data)
```

Using Pyspark native sql

```
-----

from pyspark.sql import SQLContext, Row
sqlContext = SQLContext(sc)
sqlContext.sql("set spark.sql.shuffle.partitions=10");

ordersRDD = sc.textFile("/user/cloudera/sqoop_import/orders")
ordersMap = ordersRDD.map(lambda o: o.split(","))
orders = ordersMap.map(lambda o: Row(order_id=int(o[0]), order_date=o[1], \
order_customer_id=int(o[2]), order_status=o[3]))
ordersSchema = sqlContext.inferSchema(orders)
```

```
ordersSchema.registerTempTable("orders")
```

```
orderItemsRDD = sc.textFile("/user/cloudera/sqoop_import/order_items")
```

```
orderItemsMap = orderItemsRDD.map(lambda oi: oi.split(","))
```

```
orderItems = orderItemsMap.map(lambda oi: Row(order_item_id=int(oi[0]),  
order_item_order_id=int(oi[1]), \
```

```
order_item_product_id=int(oi[2]), order_item_quantity=int(oi[3]),  
order_item_subtotal=float(oi[4]), \
```

```
order_item_product_price=float(oi[5])))
```

```
orderItemsSchema = sqlContext.inferSchema(orderItems)
```

```
orderItemsSchema.registerTempTable("order_items")
```

```
joinAggData = sqlContext.sql("select o.order_date, sum(oi.order_item_subtotal), \  
count(distinct o.order_id) from orders o join order_items oi \  
on o.order_id = oi.order_item_order_id \  
group by o.order_date order by o.order_date")
```

```
for data in joinAggData.collect():
```

```
    print(data)
```

```
=====
```

SORTING IN SQL

```
-----
```

#Global sorting and ranking

```
select * from products order by product_price desc;
```

```
select * from products order by product_price desc limit 10;
```

#By key sorting

#Using order by is not efficient, it serializes

```
select * from products order by product_category_id, product_price desc;
```

#Using distribute by sort by (to distribute sorting and scale it up)

```
select * from products distribute by product_category_id sort by product_price  
desc;
```

#By key ranking (in Hive we can use windowing/analytic functions)

```
select * from (select p.*,  
dense_rank() over (partition by product_category_id order by product_price desc)  
dr  
from products p  
distribute by product_category_id) q
```

=====

GET NUMBER OF CANCELED ORDERS IN MONTH OF JANUARY IN 2014

```
select order_status, count(1) from orders  
where from_unixtime(cast(substr(order_date, 1, 10) as int)) like '2014-01%' group  
by order_status;  
where dr <= 2 order by product_category_id, dr;
```